

Arya by Arteq AI

The AI Voice Receptionist for Hospitals & Clinics

Founder Update - What we have built, and the business behind it

Prepared June 2026 | All figures in Indian Rupees (shown as 'Rs'). | Prices exclude 18% GST.

In one sentence

Arya answers every hospital phone call in Malayalam and 5 other languages, 24x7, on every line at once - she books appointments, confirms them on WhatsApp, manages the patient queue, alerts staff for emergencies, and never puts a caller on hold.

What's inside

- 1 What we shipped - the full product, in plain language
- 2 How it works - workflow from every point of view
- 3 Pricing plans
- 4 What it cost us to build (fixed costs)
- 5 What it costs to run, per minute
- 6 Ways we can deliver the service, and their cost
- 7 Expected monthly recurring revenue (MRR)
- 8 Where we stand & what's next

1. What we shipped

Arya is live and feature-complete. Below is everything she can do today, in plain language. Each item is fully working, not a plan.

Answering & conversation

- **6 languages.** Arya speaks Malayalam, Hindi, Tamil, Telugu, Kannada and English, and switches to the caller's language automatically.
- **Always on.** She picks up instantly and handles unlimited calls at the same time - no hold music, no queue, no missed calls after hours.
- **Human-like.** She sounds natural and replies fast. Common phrases are pre-voiced and reused, so there is no robotic delay.

Appointments & the patient queue

- **Self-service booking.** Patients can book, reschedule or cancel an appointment entirely by voice. Arya knows doctors, departments, and timings.
- **Confirmation code.** Every booking gets a short confirmation code (e.g. ARYA-7K2P) the patient can quote at the desk.
- **Pay-then-token.** The patient pays the fee at the hospital; the moment staff mark it paid, Arya activates the patient's queue token and messages them the token number. No payment, no token - clean and fair.
- **Smart priority.** Emergencies and senior citizens are automatically pushed up the queue, so the most urgent patients are seen first.
- **Load balancing.** If a department has several doctors, Arya routes the patient to the least-busy one, balancing the load across the team.
- **No double-booking.** Two patients grabbing the same slot at the same instant can never double-book - the system locks the slot safely.

Messaging & staff alerts

- **WhatsApp-first.** Confirmations and reminders go out on WhatsApp, with an automatic SMS fallback if WhatsApp can't be delivered.
- **Staff alerts.** Staff get an instant alert for new bookings, cancellations, and especially emergencies, so a human is always in the loop.
- **Live transfer.** Arya can transfer a live call to a human staff member when the caller asks or when she judges it necessary.

Automatic follow-ups (runs on its own)

- **Outbound engine.** Appointment reminders before the visit, confirmation requests, missed-call callbacks, and post-visit follow-ups - all sent automatically on a schedule, with no staff effort.

Dashboard & multi-hospital

- **Owner dashboard.** A web dashboard shows call logs, bookings, costs, and live activity, protected by secure staff login.
- **Multi-tenant.** One system runs many hospitals and clinics at once, each with its own doctors, persona name, and data kept private and separate.
- **Private & secure.** Hospital data stays isolated per tenant and access is authenticated - built with patient privacy in mind.

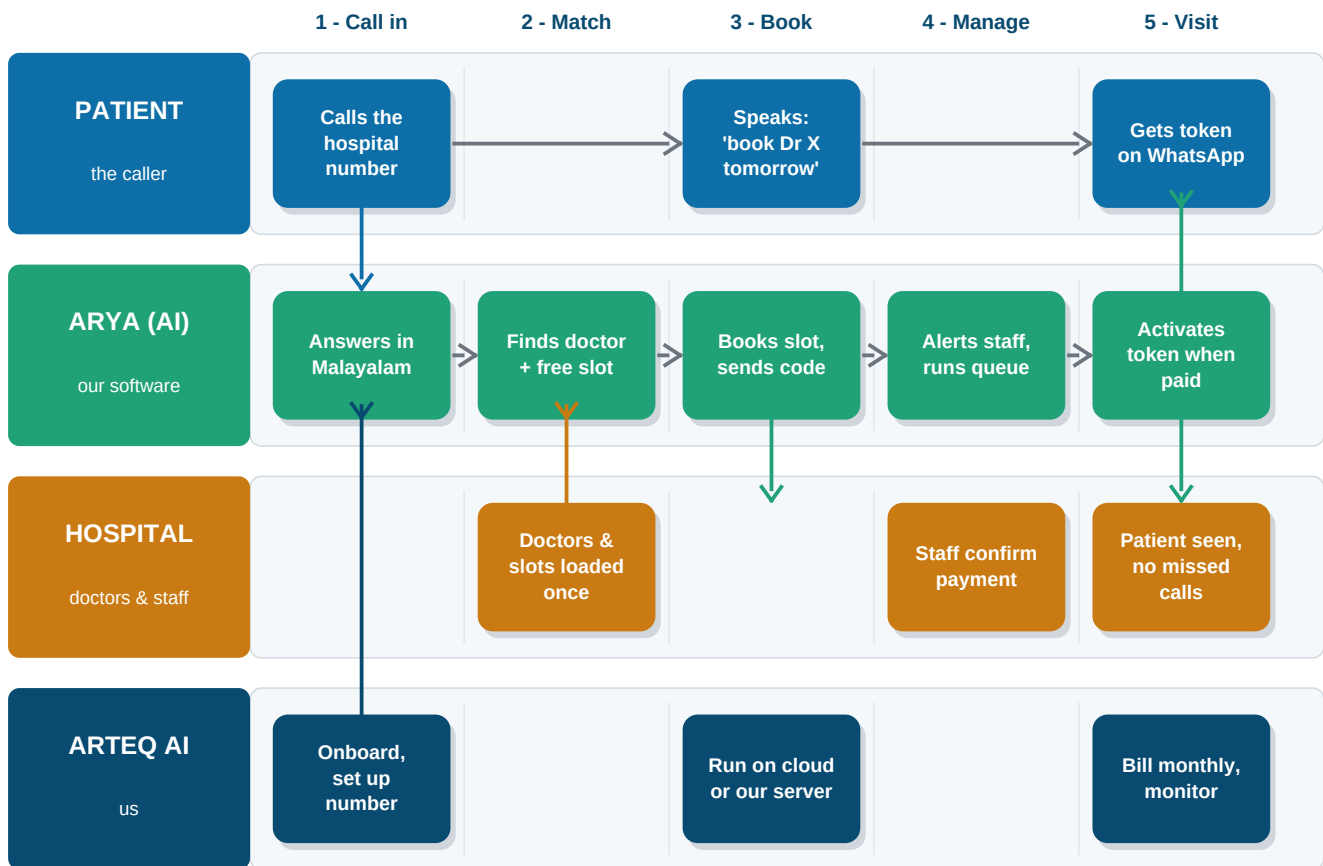
Deployment (how we ship it)

- **Two ways to run.** Arya can run on managed cloud for a fast launch, or be fully self-hosted on a

single low-cost server we control - same quality, lower cost. Both are packaged in containers, so setup is repeatable.

2. How it works - from every point of view

The same call touches four groups. Here is what each one sees and does.



Read it like this

Each row is one group; time runs left to right across the 5 stages on top. The patient just talks. Arya (green) does all the work in the middle. The hospital only loads its doctors once and confirms payments. We (Arteq) set it up, run it, and bill a simple monthly fee. Arrows show the call moving between them.

3. Pricing plans

Each plan is a monthly subscription with a block of included talk-time. Extra minutes are billed at the plan rate + Rs 0.50/min.

Plan	Included	Monthly	Rs/min	Best for
Starter	1,000 min	Rs 6,999	7.00	Single-doctor clinics
Growth	2,500 min	Rs 14,999	6.00	Small clinics
Professional	30,000 min	Rs 1,34,999	4.50	Mid-size hospitals
Enterprise	60,000 min	Rs 2,39,999	4.00	Large hospitals
Enterprise+	100,000 min	Rs 3,49,999	3.50	Multi-specialty groups

Every plan includes: 6 languages, 24x7 answering, unlimited simultaneous calls, voice booking, WhatsApp + SMS confirmations, queue tokens, priority & load balancing, live transfer, staff alerts, and the owner dashboard.

Add-ons

One-time setup (number, data load, staff training): Rs 10,000 - 25,000. Optional 14-day free trial capped at ~300 min (our cost under Rs 900). All prices exclude 18% GST.

4. What it cost us to build

These are the fixed and tooling costs to design, build and stand up Arya. Most are monthly subscriptions during the build; the engineering itself was done in-house.

Item	What it's for	Cost
Claude Code (AI dev)	Building the whole product	Rs 1,500-1,800/mo
Groq API	The AI 'brain' (reasoning)	Pay-per-use, ~Rs 0.20/min
Sarvam API	Speech-to-text + voice (Indian langs)	Pay-per-use, ~Rs 0.65/min
LiveKit	Voice call infrastructure (open-source)	Free self-host / cloud per-min
Plivo	Phone numbers + WhatsApp/SMS	~Rs 0.60/min + number rent
Supabase	Database (patient & booking data)	Free tier -> Rs ~2,000/mo
VPS (Hostinger)	Server to run Arya	Rs 799/mo (renews ~1,400)
Domain + misc	Web address, certificates	Rs ~1,000/yr

The key point for cost

Almost everything is pay-as-you-go. We don't pay for idle capacity - the big costs (Groq, Sarvam, Plivo) are only charged per minute Arya is actually on a call. Fixed monthly spend during build is roughly Rs 5,000 - 6,000, mostly tooling.

5. What it costs to run, per minute

Every minute Arya is on a call, we pay these providers. We run two setups and move customers to the cheaper self-hosted one as volume grows.

Cost component	Managed Cloud	Self-hosted VPS
Phone line (Plivo)	Rs 0.60	Rs 0.60
Voice infra (LiveKit)	Rs 1.20	~Rs 0.10
Speech understanding (Sarvam)	Rs 0.50	Rs 0.50
Speech voice (Sarvam)	Rs 0.15	Rs 0.15
AI brain (Groq)	Rs 0.20	Rs 0.20
SMS / WhatsApp + buffer	Rs 0.10	Rs 0.10
Total per minute	~Rs 2.75	~Rs 1.65

Self-hosting drops the voice-infra cost from Rs 1.20 to a few paise because one Rs 799 server is shared across many clinics (multi-tenant). Quality is identical - only where LiveKit runs changes.

6. Ways we can deliver the service

Three delivery models, same product. They trade setup speed against running cost and margin.

Option	What it is	Our cost/min	Best when
Managed Cloud	LiveKit's rented cloud; fastest launch	~Rs 2.75	First customers, pilots
Self-hosted VPS	LiveKit on our Rs 799 box, multi-tenant	~Rs 1.65	Default at scale
Dedicated VPS	One hospital's own box + standby failover	~Rs 1.85	Large hospital wanting isolation

Our recommendation

Launch new customers on Managed Cloud for a same-day start, then migrate them onto a shared self-hosted box once their volume is steady. This keeps onboarding instant while protecting our ~50% margin. Offer a dedicated box (with a standby for failover) only to large hospitals that ask for it.

Per-plan margin stays near 50% across the board: small plans run on cloud (higher cost, higher price), large plans on our own infra (lower cost).

Plan	Price	Our cost	Gross profit	Margin
Starter	Rs 6,999	~Rs 3,000	Rs 3,999	57%
Growth	Rs 14,999	~Rs 7,500	Rs 7,499	50%
Professional	Rs 1,34,999	~Rs 66,000	Rs 68,999	51%
Enterprise	Rs 2,39,999	~Rs 1,20,000	Rs 1,19,999	50%
Enterprise+	Rs 3,49,999	~Rs 1,80,000	Rs 1,69,999	49%

7. Expected monthly recurring revenue (MRR)

MRR is the predictable subscription income each month. Below are three scenarios for the first 12 months, using a realistic Kerala customer mix. 'Gross profit' uses our ~50% blended margin.

A realistic customer mix

Most clinics land in Starter/Growth; most private hospitals in Professional/Enterprise. We model a blended average revenue per customer of about Rs 22,000/month across this mix.

Scenario (by month 12)	Customers	Avg/customer	MRR	Gross profit/mo
Conservative	10	Rs 22,000	Rs 2.2 L	~Rs 1.1 L
Base case	25	Rs 22,000	Rs 5.5 L	~Rs 2.75 L
Aggressive	50	Rs 24,000	Rs 12.0 L	~Rs 6.0 L

L = lakh (Rs 1,00,000). Annualised, the base case is roughly Rs 66 lakh ARR with about Rs 33 lakh gross profit, before salaries and marketing.

Why this is reachable

- Kerala has thousands of clinics and hundreds of private hospitals - the target market is large and underserved on phone handling.
- Each closed customer is sticky: once Arya holds their doctors, queue and history, switching cost is high.
- Margin holds at ~50% even as we grow, because larger customers run on our cheaper self-hosted infrastructure.
- A single Rs 799 server can carry 20-40 small clinics, so adding customers barely adds cost.

The simple pitch

"One missed call can be one lost patient. Arya answers every call, in Malayalam, day and night, and books the appointment on the spot - for less than the cost of a single receptionist."

8. Where we stand & what's next

Done

- Full product built and working: answering, booking, WhatsApp tokens, priority queue, load balancing, alerts, dashboard, multi-hospital.
- Both delivery setups packaged (managed cloud + one-server self-host).

To go fully live

- Point the hospital's phone number/WhatsApp line at Arya (one-time setup).
- Load the first hospital's doctors, departments and timings.
- Run the first paid pilot and confirm real-world call quality and cost.

Notes: Figures are indicative, based on June 2026 provider rates - revisit if Plivo / LiveKit / Sarvam / Groq pricing changes. Add 18% GST on all customer prices. MRR scenarios assume the stated customer counts and the blended average revenue per customer; actual results depend on sales pace.